



Real-time AI risk intelligence for healthcare organizations - making AI adoption safer, not just faster.

10K+

ARTICLES DAILY

200+

ORGANIZATIONS PROTECTED

98%

ACCURACY



Stop Guessing. Start Monitoring.

Vigilara is the intelligence layer that closes the safety gap in the 40% annual healthcare annual healthcare AI growth.

REAL-TIME INTELLIGENCE

HEALTHCARE-SPECIFIC

RISK SCORING 1-10

HIPAA-AWARE

THE MISSION

Our mission: ensure every AI-assisted healthcare decision is safe, transparent, and trustworthy.

Healthcare AI grows **40% per year**. The safety infrastructure to monitor it barely exists. Vigilara is the intelligence layer that closes that gap.



Who We Serve

PRIMARY PERSONA

Dr. Madhumita Hospital CIO

- Oversees AI procurement across 12 departments
- Spends 3+ hours weekly on manual risk scanning
- Needs oversight without a dedicated dedicated analyst team

CORE PAIN POINT

"Flying blind on AI risk every week" week"

SECONDARY PERSONA

Sakshi AI Vendor

- Deploying AI tools in clinical settings
- Needs proactive monitoring to protect protect clients
- Can't afford a full compliance team

CORE PAIN POINT

"One bad AI incident could end us" us"

TERTIARY PERSONA

John Compliance Officer

- Team of 4, no deep AI expertise
- Must report AI risks to leadership quarterly
- Existing tools too technical for her team team

CORE PAIN POINT

"I need answers, not dashboards"



Our Position in the Market

"Vigilara delivers real-time, healthcare-specific AI risk intelligence — transforming hours of manual monitoring into 90 seconds of actionable seconds of actionable alerts."

WE'RE NOT

- ✗ Generic aggregator
- ✗ Checklist tool
- ✗ Research platform

WE ARE

- ✓ Healthcare-specific intelligence
- ✓ Real-time not reactive
- ✓ Built for operators

FEATURE	MANUAL SEARCH	CAIS	AI RISK EXPLORER	VIGILARA
Real-Time	No	No	No	YES
Healthcare-Specific	No	Partial	No	YES
For Non-Technical	Yes	No	No	YES
Cost	Free	Enterprise	Unknown	\$5/month



Strategy & Brand Pillars

UNIQUE VALUE PROPOSITION

"Vigilara helps healthcare organizations adopt AI safely by delivering **real-time risk intelligence** so they can move fast without flying blind."



VIGILANCE

Continuous watching so teams never have to. Real-time monitoring is our core promise.



SAFETY FIRST

Patient safety is the north star. Every risk score is anchored in clinical impact.



ACTIONABILITY

Tells you what to do, not just what happened. Insights designed for immediate response.



HEALTHCARE NATIVE

Built by someone inside the problem. Deep understanding of understanding of clinical workflows.



TRUSTWORTHY

Earns confidence before adoption. Transparent scoring and verifiable data sources.



Visual Identity System

THE LOGO



RATIONALE

The eye represents vigilance and the guardian archetype. Scalable from favicon to billboard, it balances trust (cyan) and care (pink).

COLOR & TYPOGRAPHY



HEADINGS / INTER BOLD 900

Stop Guessing.

High-impact, authoritative, and professional.

DATA / JETBRAINS MONO



Risk Score: 9.2

Used for technical metrics and risk scoring.



"Visual consistency is the foundation of user trust."



Brand Everywhere It Matters



LinkedIn Janhavi Patil, Healthcare Data Scientist



Navy → Dashboard

Cyan → Primary/Alerts

Pink → Critical Flags

Teal → Low Priority

A Inter Bold → Headings

</> Mono → Risk Scores

"The brand and the product are the same experience."

Brand Voice in Action

CONFIDENT. CLEAR. EMPOWERING. PRECISE.

PROBLEM-FOCUSED

"67% of healthcare AI systems have no real-time monitoring. The average failure takes 18 months to detect. Vigilara watches 24/7 so you never find out the hard way."

VALUE-DRIVEN

"What would you do with 3 extra hours a week? That's what you save replacing manual monitoring with Vigilara. Same intelligence. 90 seconds."

EDUCATIONAL

"Three questions every hospital should ask about their AI:
1. Is it monitored right now?
2. Who gets alerted when it fails?
3. How fast?
Vigilara answers all three —



LINKEDIN STRATEGY

Targeting CIOs and B2B buyers with authoritative insights. Voice: never alarmist, always alarmist, always empowering. 81 impressions on featured content.



SUBSTACK STRATEGY

"The Invisible Risk in Healthcare AI" — Weekly deep-dives for thought leadership and building long-term trust with clinical operators.

• STORY 1 · OVERCOMING THE MONSTER

The Invisible Risk Nobody Was Watching.

THE SETUP

Working inside healthcare AI systems I kept noticing something — models running in production, clinical decisions at scale, but **nobody watching in real time.**

RIISING ACTION

At Catalyst Inc. I watched what good intelligence infrastructure looks like. That raised the bar and **revealed the gap** more sharply than ever.

TURNING POINT

I came to Northeastern with two years of experience and a genuine desire to build something meaningful. **Not a class project** — the answer the industry needed.

THE VICTORY

Real-time scanning. Risk scores 1–10. **200+ organizations protected.** Live at vigilara.ai. The monster is no longer invisible.

BRAND CONNECTION

Monster → Vigilance
Stakes → Safety First
Answer → Native

• STORY 2 · RAGS TO RICHES

Before Vigilara. After Vigilara.

BEFORE — THE STRUGGLE

A compliance officer's Monday morning: **manually searching** Google News, reading dozens of articles, hoping nothing critical slipped through. Slow. Exhausting. Unreliable.

THE OPPORTUNITY

Healthcare AI growing 40% per year. Manual monitoring that barely worked for 10 tools would **completely break under 50.**

THE SETBACK

Every solution was too broad (all AI, not healthcare) or too technical (built for researchers). **She was still on her own.**

AFTER — THE SUCCESS

One dashboard. Risks scored 1–10. Alerts surfaced. Setup under 5 minutes. **That Monday no longer feels like a fire drill.**

BRAND CONNECTION

Struggle → Actionability
Transform → Safety First
User-first → Native

Vigilara Live Demo

HEALTHCARE AI TOPIC → ACTIONABLE RISK INTELLIGENCE IN 90 SECONDS

DEMO FLOW

- Enter Topic**
Specify any healthcare AI system or research area.
- Real-Time Scan**
Scanning 10,000+ sources via automated pipelines.
- AI Risk Scoring**
Claude API analysis across 4 clinical categories.
- Actionable Insights**
Plain-language recommendations for CIOs.

TECH STACK

- Claude API**
Healthcare-specific risk scoring and research analysis.
- n8n**
Automated workflow pipeline for continuous monitoring.
- Python**
Core data processing and intelligence layer.
- Streamlit**
Non-technical dashboard for real-time risk visualization.

LIVE ACCESS

Inputs

Sources to monitor

The Verge (tech)

Fierce Healthcare

arXiv (AI news)

Articles per source

3

15

Optional keyword filter

e.g., "clinical AI," "FDA," "risk," "patient safety"

Run Analysis

Created by: Jonathan Paul

Portfolio Contact: (add your link here in the README / app once ready)

Medical AI Risk Intelligence

An assignment AI workflow wrapped in a friendly interface. This tool scans recent healthcare / AI news and uses Claude to score medical AI risk, categorize issues, and surface key safety insights.

About this tool

- What it does: Fetches recent articles from selected tech and healthcare sources and asks Claude to score each one for medical AI risk.
- Who it's for: Product managers, policy students, and healthcare innovators who want a quick view of emerging AI risk signals.
- Tech stack: Streamlit, Python, Claude API, RSS feeds (The Verge, Fierce Healthcare, arXiv cs.AI).

Analyzing article 1/10: Embraer's self-heating smart mug is more than 500 off ahead of Mother's Day

Analyzing article 2/10: Watch Sony's elite ping-pong robot beat top-ranked players

Analyzing article 3/10: Anthropic's Mythos rollout has missed America's cybersecurity agency

Analyzing article 4/10: Microsoft teases mysterious Discord and Xbox Game Pass partnership

Analyzing article 5/10: Google Meet will take AI notes for in-person meetings too

Analyzing article 6/10: Community Health Systems attributes Q1 volumes stumble to consumers' macroeconomic fears, payers' prior auth denials

Analyzing article 7/10: Why UnitedHealth Group is betting big on doulos

Analyzing article 8/10: Oak HC/VT backs Courier Health's S50M series B to build out AI for the biopharma patient experience

Analyzing article 9/10: Amgen Health secures \$10M in series A funding, unveils AI native denial management solution

Analyzing article 10/10: Elevance Health raises 2026 outlook as it posts \$1.8B in Q1 profit

Risk Overview

LOW RISK: 5/10

3. MrBeast just bought a banking app

Source: (thetrump) | Published: 2/9/2026 | URL: <https://www.thetrump.com/news/37080/mrbeast-owns...>

market_signal

health_tech

Key Insights:

- Non-healthcare entity acquiring financial services targeting young demographics
- Potential indirect impact on healthcare AI through financial wellness platform expansion
- Market consolidation trend of content creators entering traditional service sectors

AI Safety Relevance: Minimal direct relevance to medical AI safety. This is a financial services acquisition with no apparent healthcare AI components or immediate patient safety implications.

LOW RISK: 5/10

5. Toyota made a game engine

Source: (thetrump) | Published: 2/9/2026 | URL: <https://www.thetrump.com/news/37395/toyota-built...>

market_signal

health_tech

Key Insights:

- Automotive game engine development may influence healthcare AI interface design patterns
- Cross-industry technology transfer from gaming to safety-critical systems sets precedents for medical AI
- Enhanced user interface technologies in vehicles could inform patient-facing medical device interfaces

AI Safety Relevance: While not directly medical AI, this represents technology trends in safety-critical human-machine interfaces that could influence healthcare AI system design, particularly in medical devices requiring intuitive user interaction.

LOW RISK: 2/10

4. FCC accused of withholding DOGE information 'in bad faith'

Source: (thetrump) | Published: 2/9/2026 | URL: <https://www.thetrump.com/news/37298/fcc-doge-f...>

market_signal

health_tech

Key Insights:

- Article focuses on FCC-DOGE transparency dispute rather than medical AI safety
- No direct healthcare AI system or patient safety concerns mentioned
- Regulatory transparency issues may indirectly affect health tech oversight

AI Safety Relevance: Minimal relevance - article concerns government efficiency efforts at FCC with no direct connection to medical AI systems, patient care, or healthcare technology safety

Medical AI Risk Intelligence Report

Assignment 4: Healthcare AI Safety Analysis

Generated: 2/9/2026, 10:18:33 PM

10

ARTICLES ANALYZED

2.70

AVG RISK SCORE

2

HIGH RISK

2

ACTIONABLE ITEMS

\$0.0347

PROCESSING COST

4,933

TOTAL TOKENS

Risk Distribution

2

High Risk (7-10)

0

Medium Risk (4-6)

8

Low Risk (1-3)

3-Slide Brand Pitch Deck

PITCH SLIDE 1 — THE PROBLEM

The \$50B Gap Nobody Is Watching

Healthcare AI is growing 40% per year. The safety infrastructure to monitor it barely exists. 67% of systems operate with no real-time oversight. Critical failures take 18 months to detect on average. This is the problem Vigilara was built to solve.

\$50B+ **67%** **18mo**
Market Size Unmonitored Detection Time

PITCH SLIDE 2 — THE SOLUTION

Real-Time Intelligence. Healthcare-Specific.

Vigilara scans 10,000+ articles daily using Claude API, scores risks 1–10 across four healthcare domains, and delivers actionable alerts in under 90 seconds. Built on n8n automation and Streamlit — no technical expertise required to use it.

10K+ **90sec** **98%**
Articles/Day Alert Time Accuracy

PITCH SLIDE 3 — THE ASK

Start Free. Scale When Ready.

Free plan for exploration. Pro at \$99/month for growing organizations monitoring up to 25 AI systems. Enterprise at \$499/month for unlimited systems, custom models, and dedicated support. 200+ healthcare organizations already protected.

\$0 **\$99** **200+**
Free Forever Pro/Month Orgs Protected

Every slide built using Vigilara's complete brand system — navy, cyan, pink, Inter typography, eye icon

vigilara.ai



Thank You.

10K+

ARTICLES DAILY

200+

PROTECTED

98%

ACCURACY

Free

TO START

Real technology. Real users. Real impact on how healthcare organizations adopt AI safely.

vigilara.ai

WEBSITE

github.com/janhavi502

GITHUB

patil.janhavi@northeastern.edu

EMAIL

ANY QUESTIONS?

JANHAVI PATIL | MS INFORMATION SYSTEMS | NORTHEASTERN UNIVERSITY

