

Your taste, on autopilot.

The only research tool that holds an opinion on your behalf — and gets provably better at it the more you use it.

Not a content tool. A verdict tool.



Made with **GAMMA**

The problem isn't information. It's **attention**.

840+ papers/week on arxiv cs.AI · **4,300+** papers/month across AI sources · No tool holds opinions

PRIMARY USER PERSONA

Maya — AI PM

Series C startup · 28–35

8:45 AM. 47 unread Slack messages.
Eight are links. Her CEO just asked about a paper from four days ago.

She has the information. She doesn't have a position.

- Pain: Information volume vs. decision velocity
- Budget: \$50–200/mo for tooling

SECONDARY USER PERSONA


Jordan — Sr. Research Analyst

AI-focused VC · 30–42

Covers frontier AI for the investment committee. Makes fewer, higher-stakes calls on fundraise-worthy companies.

Bespoke research is expensive. Generic AI tools don't hold opinions.

- Pain: Needs a tool that decides, not summarises
- Budget: \$500–2,000/mo team tier



A different category.

Dimension	Chatbots	Aggregators	AI Branding Tools	SignalScout
Mode	Ask → answer	Aggregate → summarise	Generate → produce	Evaluate → decide
Output	Content	More content	Even more content	Verdict
Personalisation	Same for everyone	Same for everyone	Same for everyone	Diverges per user
Memory	None	None	None	Learns you. Compounds.
Proof of learning	Static	Static	Static	14× random after 500 decisions

The market is flooded with tools that outsource writing.
SignalScout outsources reading.

Four pillars. Each reinforces the others.

1

Decision over content.

Produce less. Decide more. Compete on velocity, not volume.

2

Accountability of taste.

Show the reasoning. Show the track record. No black boxes.

3


Personal over generic.

Your SignalScout is mathematically different from anyone else's. Architecture guarantees divergence per user.

4


Silence over noise.

The UI is the thesis. Reduction is the product.
Bloomberg terminal, not ChatGPT.



The brand system.

PALETTE



TYPOGRAPHY

Space Grotesk

Headers — sharp, modern, no decoration.

Inter

Body copy — the most readable sans-serif.

JetBrains Mono

+444 reward · 14x · 28ms — numbers earn monospace.




Three fonts. Four Logo Iterations. Four colors. One rule: *numbers earn monospace.*

Three minds argue. A **bandit** learns you.

SKEPTIC
Judges rigour.
 Penalises buzzwords. Rewards reproducibility.
 Baseline: **40 / 100**

SCOUT
Hunts novelty.
 Rewards weak signals from frontier labs.
 Baseline: **50 / 100**

OPERATOR
Judges shippable.
 Rewards deployable techniques this quarter.
 Baseline: **45 / 100**



Three evaluators score each paper in parallel — under 2 sec wall-clock. A contextual multi-armed bandit picks which evaluator to trust for this specific user on this specific item.
15-dimension feature vector. Persistent memory across sessions. Feedback updates the posterior via rank-1 Bayesian update. Sublinear regret, verified in CI tests.

30
Decisions
 Personalisation starts working

14×
Better
 than random baseline

500
Decisions
 Bandit knows you better than baseline

+444
Reward gap
 vs. random baseline

THREE PARALLEL EVALUATORS

Skeptic, Scout, Operator
score every paper in parallel
— under 2 sec wall-clock.

PERSISTENT TASTE MEMORY

Describes you back to
yourself. Compounds across
sessions.

CONTEXTUAL BANDIT

Picks which evaluator to trust
per user, per item. Learns,
doesn't guess.

DAILY EMAIL BRIEF

Ranked. Reasoned. Personal.
Ink Black header band. Logo
lockup.

METHODOLOGY

The brand's four pillars aren't decoration — they're architecture constraints.

"Decision over content" → one ranked brief, not an endless feed.

"Accountability of taste" → every score has a visible rationale.


"Personal over generic" → bandit state is persistent per user.

"Silence over noise" → UI reduces to ranked verdicts, nothing else.

Brand, applied to every surface.

LIVE AGENT WEBSITE


URL: signalscout.streamlit.app



- Logo lockup in sidebar — consistent across all views
- Mood cards per evaluator — brand colours as functional UI
- Tagline + promise in page header
- JetBrains Mono on every number, score, and latency

MARKETING WEBSITE

URL: signalscout-831047036020.us-west1.run.app



BRAND IN PRODUCT

The brand isn't skinned on top. The brand is how the product **behaves**.

- Evaluators glow in their own colour when they win the bandit's pick.
- Every number — scores, latencies, reward — renders in JetBrains Mono.
- The email brief opens with an Ink Black header band and the logo lockup.





PLATFORM STRATEGY

LinkedIn — Primary

Where AI PMs work in 2026. Technical substance performs.

Substack/ Medium — Secondary


Reaches VC and analyst personas. Shorter, sharper.



Karpathy tweets. Everyone quotes him.

Almost nobody has read the paper he's quoting.

The AI community runs on secondhand summaries.



By the time a paper reached me, it had been through six layers of compression and two rounds of wishful thinking.

I don't want to read Karpathy's read of the paper.

I want to know if the paper is worth my read.

ORDINARY WORLD

Reading influencer threads about papers I never opened. Mistaking proximity to smart people for having a take.

INCITING INCIDENT

Realising my "opinions" were six hops from the source. That my taste was borrowed, not built.

TRIALS → TRANSFORMATION

Six evaluator prompts. Four bandit rewrites. Three months arguing with my own tool.

Now the tool has the take. I just approve or reject it.

Brand values reinforced:

Decision over content · Accountability of taste · Silence over noise

Products that come from scar tissue feel **inevitable**. Not invented.

The system is learning. Here's the receipt.

BEFORE SIGNALSCOUT


8:45 AM. 47 unread Slack links. Eight papers to read. Two hours to skim. Still no position on Constitutional AI 2.0 when the CEO asks.

AFTER SIGNALSCOUT

8:45 AM. One email in her inbox. "Your brief for April 24." Seven papers. Ranked. Reasoned. Personal. She walks into standup with a point of view.

Framework: Before / After transformation • Brand value: Decision over content

— Random Baseline — SignalScout



500

Decisions

Tested in simulation

+444

Reward Gap

Cumulative gap vs. random

14×

Better

Than random baseline

"You reward Scout 2.3× more on items tagged 'agents.' Sceptic's picks correlate with your 👍 at 64%."

— SignalScout describing a user to themselves

Your taste, on autopilot.

Stop reading everything. Start knowing what matters.

signalscoutai.streamlit.app